

# Kubernetes

## Notes

Issues with local ARM based cluster - non comptabile images

Runing on AKS, with one slight issue in [probes\\_http](#)

## 1. Simple pods operations:

No pods available after querying

```
kubectl get pods
No resources found in default namespace.
```

While using --all-namespaces we see pods(21) under the kube-system namespace. These pods are made and owned by the cluster - basically services that make using k8s possible

```
kubectl get pods --all-namespaces
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	ama-logs-qxsfd	2/2	Running	0	17m
kube-system	ama-logs-rs-6d99f5c8dd-d4476	1/1	Running	0	11h
kube-system	azure-ip-masq-agent-l6h57	1/1	Running	0	17m
kube-system	cloud-node-manager-9gsvr	1/1	Running	0	17m
kube-system	coredns-59b6bf8b4f-4ptgd	1/1	Running	0	11h
kube-system	coredns-59b6bf8b4f-ntb9n	1/1	Running	0	11h
kube-system	coredns-autoscaler-5f9cb57949-sqfqz	1/1	Running	0	11h
kube-system	csi-azuredisk-node-7hr9k	3/3	Running	0	17m
kube-system	csi-azurefile-node-4c577	3/3	Running	0	17m
kube-system	connectivity-agent-75f989679d-rbbbg	1/1	Running	0	11h
kube-system	connectivity-agent-75f989679d-wtzdg	1/1	Running	0	11h
kube-system	kube-proxy-5qt2v	1/1	Running	0	17m
kube-system	metrics-server-7dd74d8758-j4vrz	2/2	Running	0	16m
kube-system	metrics-server-7dd74d8758-rfk6v	2/2	Running	0	16m

Afte executing *kubectl run nginx --image=nginx* and querying the pod is runing and well.

```
kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx	1/1	Running	0	12s

Running *logs* gives information about the pod - version, process status etc

```
kubectl logs nginx
```

```
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
2023/04/05 05:49:00 [notice] 1#1: using the "epoll" event method
2023/04/05 05:49:00 [notice] 1#1: nginx/1.23.4
2023/04/05 05:49:00 [notice] 1#1: built by gcc 10.2.1 20210110 (Debian 10.2.1-6)
2023/04/05 05:49:00 [notice] 1#1: OS: Linux 5.4.0-1104-azure
2023/04/05 05:49:00 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1048576:1048576
2023/04/05 05:49:00 [notice] 1#1: start worker processes
2023/04/05 05:49:00 [notice] 1#1: start worker process 29
2023/04/05 05:49:00 [notice] 1#1: start worker process 30
```

Using *-o wide* gives us more information including the node where the pod is running

```
kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
nginx	1/1	Running	0	33m	10.224.0.80	aks-agentpool-41905252-vmss000003	<none>	<none>

Using *describe* gives us a detailed state of a pod

```
kubectl describe pod nginx
Name:          nginx
Namespace:     default
Priority:       0
Service Account: default
Node:          aks-agentpool-41905252-vmss000003/10.224.0.4
Start Time:    Wed, 05 Apr 2023 08:48:56 +0300
Labels:        run=nginx
Annotations:    <none>
Status:        Running
IP:            10.224.0.80
IPs:
```

Container Image is pulled from a microsoft repository or better know as Microsoft Artifact Registry

```
kubectl describe pod coredns-59b6bf8b4f-4ptgd -n kube-system
Name:          coredns-59b6bf8b4f-4ptgd
Namespace:     kube-system
Priority:       2000001000
Priority Class Name: system-node-critical
Service Account: coredns
Node:          aks-agentpool-41905252-vmss000003/10.224.0.4
Start Time:    Wed, 05 Apr 2023 08:23:38 +0300
Labels:        k8s-app=kube-dns
                kubernetes.io/cluster-service=true
                pod-template-hash=59b6bf8b4f
                version=v20
Annotations:    prometheus.io/port: 9153
Status:        Running
IP:            10.224.0.22
IPs:
  IP:          10.224.0.22
Controlled By: ReplicaSet/coredns-59b6bf8b4f
Containers:
  coredns:
    Container ID: containerd://cdf5833b7f3f0ae28efd85f195604337b2f92b1377af1d1aba5b0e95d22f949a
    Image:        mcr.microsoft.com/oss/kubernetes/coredns:v1.9.3
    Image ID:     sha256:c3819500b42300c8eed00da07da000002aa92727a908589804cadb445f6df72d6
    Ports:        53/UDP, 53/TCP, 9153/TCP
    Host Ports:   0/UDP, 0/TCP, 0/TCP
```

Logs from the metrics container - Errors aren't errors but information(go figure MS logging system)

```
kubectll logs metrics-server-7dd74d8758-j4vrz -n kube-system
Defaulted container "metrics-server-vpa" out of: metrics-server-vpa, metrics-server
ERROR: logging before flag.Parse: I0405 05:24:14.446140      1 pod_nanny.go:68] Invoked by [/pod_nanny --config-dir=/etc/config --cpu=44m --extra-cpu=0.5m --memory=51Mi --extra-memory=4Mi --poll-period=180000 --threshold=5 --deployment=metrics-server --container=metrics-server]
ERROR: logging before flag.Parse: I0405 05:24:14.446180      1 pod_nanny.go:69] Version: 1.8.14
ERROR: logging before flag.Parse: I0405 05:24:14.446209      1 pod_nanny.go:85] Watching namespace: kube-system, pod: metrics-server-7dd74d8758-j4vrz, container: metrics-server.
ERROR: logging before flag.Parse: I0405 05:24:14.446214      1 pod_nanny.go:86] storage: MISSING, extra_storage: 0Gi
ERROR: logging before flag.Parse: I0405 05:24:14.447149      1 pod_nanny.go:189] Failed to read data from config file "/etc/config/NannyConfiguration": open /etc/config/NannyConfiguration : no such file or directory, using default parameters
ERROR: logging before flag.Parse: I0405 05:24:14.447168      1 pod_nanny.go:116] cpu: 44m, extra_cpu: 0.5m, memory: 51Mi, extra_memory: 4Mi
ERROR: logging before flag.Parse: I0405 05:24:14.447179      1 pod_nanny.go:145] Resources: [{Base:{i:{value:44 scale:-3} d:{Dec:<nil>} s:44m Format:DecimalSI} ExtraPerNode:{i:{value:5 scale:-4} d:{Dec:<nil>} s: Format:DecimalSI} Name:cpu} {Base:{i:{value:53477376 scale:0} d:{Dec:<nil>} s:51Mi Format:BinarySI} ExtraPerNode:{i:{value:4194304 scale:0} d:{Dec:<nil>} s:4Mi Format:BinarySI} Name:memory}]
```

ToDo:

- What is a pod\_nanny

## 2. Working with pod manifest files:

For the sake of being brief I'll compare both files - wrong and correct versions.

- apiVersion is wrong for pods should be v1 not v11
- Formatting isn't according to yaml standard
- it's spec, not specs - typo
- image redis123 doesn't exist(we get a pull image error), we should check the image repo(Docker hub) for proper naming/version

```
- apiVersion:v11      → 1+ apiVersion: v1
- kind:pod            2+ kind: Pod
  metadata:          3  metadata:
- name:static-web     → 4+   name: static-web
- labels:             5+   labels:
- role:myrole         6+   role: myrole
- specs:              7+ spec:
- containers:         8+   containers:
-   name:redis        9+   name: redis
- image:redis123     10+   image: redis:latest
                    11+   ports:
                    12+   - containerPort: 6379
```

First run and we get a parsing error

```
~/Homework/Homework-13/manifets master !1 72
kubectll apply -f r3dis.yaml
error: error parsing r3dis.yaml: error converting YAML to JSON: yaml: line 3: mapping values are not allowed in this context
```

After the file has been corrected we get another error this time because of the image name

```
@ ~/Homework/Homework-13/manifets master !1 ?2
kubectll get pods
NAME          READY   STATUS    RESTARTS   AGE
static-web    1/1     Running   0           12s

└─ kubectll get pods
NAME          READY   STATUS    RESTARTS   AGE
static-web    0/1     ErrImagePull 0           10s
```

And here we have a healthy pod with a container.

```
@ ~/Homework/Homework-13/manifets master !1 ?2
kubectll get pods
NAME          READY   STATUS    RESTARTS   AGE
static-web    1/1     Running   0           12s
```

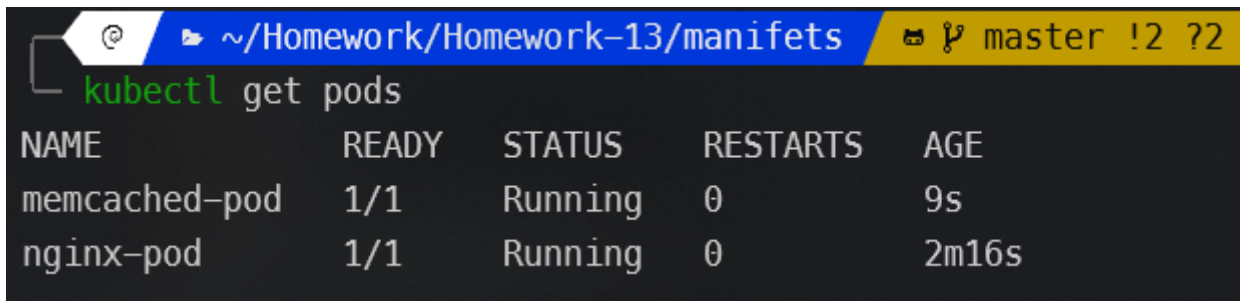
nginx manifest

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  labels:
    app: frontend
spec:
  containers:
    - name: nginx-container
      image: nginx
      ports:
        - containerPort: 80
```

## Memecache manifest

```
apiVersion: v1
kind: Pod
metadata:
  name: memcached-pod
  labels:
    app: web
spec:
  containers:
    - name: memcached-container
      image: memcached
      ports:
        - containerPort: 11211
      resources:
        requests:
          cpu: 350m
          memory: 150Mi
        limits:
          cpu: 500m
          memory: 250Mi
      restartPolicy: Never
```

Both nginx and memcache running



```
kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
memcached-pod	1/1	Running	0	9s
nginx-pod	1/1	Running	0	2m16s

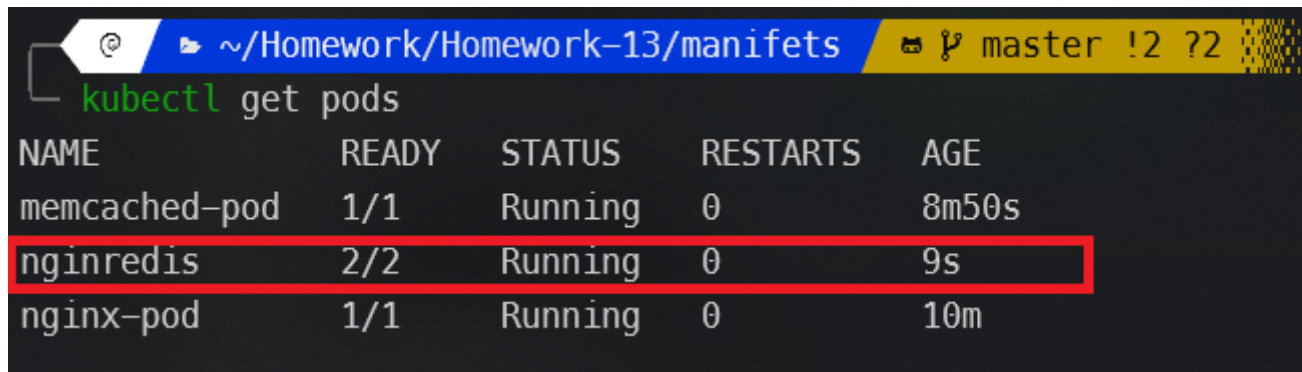
## 3. Multicontainer pods

Definition of our multicontainer pod



```
apiVersion: v1
kind: Pod
metadata:
  name: nginredis
  labels:
    app: frontend-backend
spec:
  containers:
    - name: nginx-container
      image: nginx:1.22.1
      ports:
        - containerPort: 80
    - name: redis-container
      image: redis:latest
      ports:
        - containerPort: 6379
```

The pod running with two containers inside



```
~/Homework/Homework-13/manifets master !2 ?2
kubect! get pods
```

NAME	READY	STATUS	RESTARTS	AGE
memcached-pod	1/1	Running	0	8m50s
nginredis	2/2	Running	0	9s
nginx-pod	1/1	Running	0	10m

Both containers are pulled from their respective image addresses

```
Containers:
  nginx-container:
    Container ID:   containerd://723b00ce71de88010272b6322debf7bb1e82b1d320b64872dfcf46728e62a94e
    Image:          nginx:1.22.1
    Image ID:       docker.io/library/nginx@sha256:fc5f5fb7574755c306aaf88456ebf0b006420a184d52b923d2f0197108f6b7
    Port:          80/TCP
    Host Port:      0/TCP
    State:          Running
      Started:      Wed, 05 Apr 2023 11:40:25 +0300
    Ready:          True
    Restart Count:  0
    Environment:    <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-dkt9g (ro)
  redis-container:
    Container ID:   containerd://7467493cc297e9da70c8a2df9a8f700a6648a69aeb7d070b4624b648e86bf104
    Image:          redis:latest
    Image ID:       docker.io/library/redis@sha256:7b83a0167532d4320a87246a815a134e19e31504d85e8e55f0bb5bb9edf70448
    Port:          6379/TCP
    Host Port:      0/TCP
    State:          Running
      Started:      Wed, 05 Apr 2023 11:40:25 +0300
    Ready:          True
    Restart Count:  0
    Environment:    <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-dkt9g (ro)
```

## 4. Probes

- Liveness probe

Edited yaml file for probes\_exec



```

apiVersion: v1
kind: Pod
metadata:
  labels:
    test: liveness
  name: liveness-exec
spec:
  containers:
    - name: liveness
      image: k8s.gcr.io/busybox
      args:
        - /bin/sh
        - -c
        - touch /tmp/healthy; sleep 30; rm -rf /tmp/healthy; sleep 600
      livenessProbe:
        exec:
          command:
            - cat
            - /tmp/healthy
          initialDelaySeconds: 5
          periodSeconds: 5

```

Running the describe command - no failure yet

node/k8s-1.10.1/ami-ami-8a3b2c06/execute -p EX2519 -t01-5000

Events:				
Type	Reason	Age	From	Message
Normal	Scheduled	5s	default-scheduler	Successfully assigned default/liveness-exec to aks-agentpool-41905252-vmss000000
Normal	Pulling	4s	kubelet	Pulling image "busybox"
Normal	Pulled	4s	kubelet	Successfully pulled image "busybox" in 156.36596ms
Normal	Created	4s	kubelet	Created container liveness
Normal	Started	4s	kubelet	Started container liveness

Getting the failure after 35 seconds.

node/k8s-1.10.1/ami-ami-8a3b2c06/execute -p EX2519 -t01-5000

Events:				
Type	Reason	Age	From	Message
Normal	Scheduled	38s	default-scheduler	Successfully assigned default/liveness-exec to aks-agentpool-41905252-vmss000007
Normal	Pulling	37s	kubelet	Pulling image "busybox"
Normal	Pulled	37s	kubelet	Successfully pulled image "busybox" in 156.36596ms
Normal	Created	37s	kubelet	Created container liveness
Normal	Started	37s	kubelet	Started container liveness
Warning	Unhealthy	2s	kubelet	Liveness probe failed: cat: can't open '/tmp/healthy': No such file or directory

Pod restarts have icrementated

```
~ /Homework/Homework-13/manifets master wip !2 ?1
k get pods -w
```

NAME	READY	STATUS	RESTARTS	AGE
liveness-exec	1/1	Running	1 (26s ago)	102s
liveness-exec	1/1	Running	2 (1s ago)	2m32s
liveness-exec	1/1	Running	3 (1s ago)	3m47s

- **liveness\_http**

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    test: liveness
  name: liveness-http
spec:
  containers:
    - name: liveness
      image: k8s.gcr.io/liveness
      args:
        - /server
      livenessProbe:
        httpGet:
          path: /healthz
          port: 8080
          httpHeaders:
            - name: Custom-Header
              value: Awesome
        initialDelaySeconds: 3
        periodSeconds: 3
```

## healthz is implemented in Go

Some issues related to running the healthz, might be related to incorrectly configured AKS

Events:				
Type	Reason	Age	From	Message
Normal	Scheduled	2m7s	default-scheduler	Successfully assigned default/liveness-http to aks-agentpool-41905252-vmss000009
Normal	Pulled	2m7s	kubelet	Successfully pulled image "k8s.gcr.io/liveness" in 282.344487ms
Normal	Pulled	109s	kubelet	Successfully pulled image "k8s.gcr.io/liveness" in 254.231974ms
Normal	Created	91s (x3 over 2m7s)	kubelet	Created container liveness
Normal	Started	91s (x3 over 2m7s)	kubelet	Started container liveness
Normal	Pulled	91s	kubelet	Successfully pulled image "k8s.gcr.io/liveness" in 259.054744ms
Normal	Pulling	73s (x4 over 2m7s)	kubelet	Pulling image "k8s.gcr.io/liveness"
Warning	Unhealthy	73s (x9 over 115s)	kubelet	Liveness probe failed: HTTP probe failed with statuscode: 500
Normal	Killing	73s (x3 over 109s)	kubelet	Container liveness failed liveness probe, will be restarted

ToDo:

- fix this

- probes\_tcp

```
apiVersion: v1
kind: Pod
metadata:
  name: liveness-tcp
labels:
  app: goproxy
spec:
  containers:
  - name: goproxy
    image: k8s.gcr.io/goproxy:0.1
    ports:
    - containerPort: 8080
    livenessProbe:
      tcpSocket:
        port: 9999 # 8080 is a valid port
      initialDelaySeconds: 15
      periodSeconds: 20
```

liveness-tcp is up and running

```
~/Homework/Homework-13/manifets master wip !2 74
k get pods -A
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
default	liveness-http	0/1	CrashLoopBackOff	7 (48s ago)	7m52s
default	liveness-tcp	1/1	Running	0	29s

Probe failure due to using port 9999

```
node.kubernetes.io/unreachable:NoExecute op-Exists for 300s
```

Type	Reason	Age	From	Message
Normal	Scheduled	37s	default-scheduler	Successfully assigned default/liveness-tcp to aks-agentpool-41905252-vmss000009
Normal	Pulled	37s	kubelet	Container image "k8s.gcr.io/goproxy:0.1" already present on machine
Normal	Created	37s	kubelet	Created container goproxy
Normal	Started	37s	kubelet	Started container goproxy
Warning	Unhealthy	17s	kubelet	Liveness probe failed: dial tcp 10.224.0.10:9999: connect: connection refused

- readiness\_http

```

apiVersion: v1
kind: Pod
metadata:
  name: readiness-http
  labels:
    app: test
spec:
  containers:
  - name: nginx
    image: nginx
    ports:
    - containerPort: 80
    readinessProbe:
      httpGet:
        path: /
        port: 80
        httpHeaders:
        - name: Host
          value: myapplication1.com
      initialDelaySeconds: 1
      periodSeconds: 2
      timeoutSeconds: 1
      successThreshold: 1
      failureThreshold:

```

Pod is up and running

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
default	liveness-http	0/1	CrashLoopBackOff	9 (2m30s ago)	15m
default	liveness-tcp	1/1	Running	6 (115s ago)	7m56s
default	readiness-http	1/1	Running	0	4s

No issues when running the *describe* command.

Events:					
Type	Reason	Age	From	Message	
Normal	Scheduled	91s	default-scheduler	Successfully assigned default/readiness-http to aks-agentpool-41905252-vmss000009	
Normal	Pulling	91s	kubelet	Pulling image "nginx"	
Normal	Pulled	91s	kubelet	Successfully pulled image "nginx" in 244.77341ms	
Normal	Created	90s	kubelet	Created container nginx	
Normal	Started	90s	kubelet	Started container nginx	

The pod is running but it's not ready

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
default	liveness-http	0/1	CrashLoopBackOff	9 (39s ago)	13m
default	liveness-tcp	0/1	CrashLoopBackOff	5 (4s ago)	6m5s
default	readiness-http	0/1	Running	0	32s

Due to the change the probe fails.

Events:				
Type	Reason	Age	From	Message
Normal	Scheduled	72s	default-scheduler	Successfully assigned default/readiness-http to aks-agentpool-41905252-vmss000009
Normal	Pulling	72s	kubelet	Pulling image "nginx"
Normal	Pulled	72s	kubelet	Successfully pulled image "nginx" in 147.784144ms
Normal	Created	72s	kubelet	Created container nginx
Normal	Started	72s	kubelet	Started container nginx
Warning	Unhealthy	34s (x21 over 71s)	kubelet	Readiness probe failed: Get "http://10.224.0.23:81/": dial tcp 10.224.0.23:81: connect: connection refused